

## Section 1

## Execution

The code can be executed with `./run.sh <num_time_steps>`

## Section 2

## Explanation of the code

The main code is in `src.c` which performs one run of message parsing by 3 different ways. Following is the general schema for each run:

1. Initialize data as a dynamic (but contiguous) block of size  $N*N$  with random **double** in the range  $[0, 1]$ .
2. Repeat the following steps `num_time_steps` times:
  - Compute the stencil for the data.
  - Send/Recv in all possible directions by a specific methodology.
  - Compute the final halo (only the boundaries).
3. Free the data, and related data structures.

---

**Schema 1: General Structure of the code**


---

```

1 double stime = MPI_Wtime();
2 for(int t=0; t<num_time_steps; t++){
3     // Perform stencil computation
4     compute_stencil(data, side_len);
5
6     // Send to possible directions
7     if(can_transfer('u', my_rank, cluster_len, p)){
8         // If it is possible to transfer upwards, send.
9     }
10    if(can_transfer('d', my_rank, cluster_len, p)){
11        // If it is possible to transfer downwards, send.
12    }
13    if(can_transfer('l', my_rank, cluster_len, p)){
14        // If it is possible to transfer leftwards, send.
15    }
16    if(can_transfer('r', my_rank, cluster_len, p)){
17        /// If it is possible to transfer rightwards, send.
18    }
19
20    // Recieve from possible directions
21    if(can_transfer('u', my_rank, cluster_len, p)){
22        // If it is possible to transfer upwards, receive.
23    }
24    if(can_transfer('d', my_rank, cluster_len, p)){
25        // If it is possible to transfer downwards, receive.
26    }
27    if(can_transfer('l', my_rank, cluster_len, p)){
28        // If it is possible to transfer leftwards, receive.
29    }

```

```

30     if(can_transfer('r', my_rank, cluster_len, p)){
31         // If it is possible to transfer rightwards, receive.
32     }
33
34     // Get the final averages for time t
35     compute_halo(data, recv_data, side_len, my_rank, cluster_len, p);
36 }
37 double ftime = MPI_Wtime();
38
39 double time = ftime - stime;
40 MPI_Reduce (&time, &maxTime, 1, MPI_DOUBLE, MPI_MAX, 0, MPI_COMM_WORLD);
41
42 // print only by rank 0
43 if (!my_rank) printf ("%lf\n", maxTime);

```

Following the above schema, the three strategies for message send/receive can be implemented inside each of the **if** condition. Following are three excerpts from our code for the three given strategies.

---

### Part 1: Single MPI\_Send/Recv

---

```

1 // sending upwards
2 if(can_transfer('u', my_rank, cluster_len, p)){
3     int target = my_rank - cluster_len;
4     for(int i=0; i<side_len; i++){
5         MPI_Send(&data[0*side_len + i], 1, MPI_DOUBLE, target, i, MPI_COMM_WORLD);
6     }
7 }
8
9 // recv from below
10 if(can_transfer('d', my_rank, cluster_len, p)){
11     int source = my_rank + cluster_len;
12     MPI_Status st;
13     for(int i=0; i<side_len; i++){
14         MPI_Recv(&recv_data[1][i], 1, MPI_DOUBLE, source, i, MPI_COMM_WORLD, &st);
15     }
16 }

```

Here we need to make side\_len number of MPI\_Send() MPI\_Recv() calls for transfer in a single direction.

---

### Part 2: MPI\_Send/Recv using MPI\_Pack/Unpack

---

```

1 // sending upwards
2 if(can_transfer('u', my_rank, cluster_len, p)){
3     int target = my_rank - cluster_len, upos = 0;
4     for(int j=0; j<side_len; j++){
5         MPI_Pack (&data[0*side_len + j], 1, MPI_DOUBLE, s_ubuf, side_len*8, &upos, MPI_COMM_WORLD);
6     }
7     MPI_Send(s_ubuf, side_len*8, MPI_PACKED, target, 1, MPI_COMM_WORLD);
8 }
9
10 // recv from below
11 if(can_transfer('d', my_rank, cluster_len, p)){
12     int source = my_rank + cluster_len, curr_pos = 0;
13     MPI_Status st;
14     MPI_Recv(r_dbuf, side_len*8, MPI_PACKED, source, 1, MPI_COMM_WORLD, &st);

```

```

15 MPI_Unpack(r_dbuf, side_len*8, &curr_pos, recv_data[1], side_len, MPI_DOUBLE, MPI_COMM_WORLD);
16 }

```

Here we need to make a single MPI\_Send() MPI\_Recv() call for transfer in a single direction with side\_len number of MPI\_Pack()/MPI\_Unpack().

---

### Part 3: MPI\_Send/Recv using MPI\_Datatype Vector

---

```

1 // define datatypes for row and col
2 MPI_Datatype row_vector, col_vector;
3 MPI_Type_vector (side_len, 1, 1, MPI_DOUBLE, &row_vector);
4 MPI_Type_commit (&row_vector);
5 MPI_Type_vector (side_len, 1, side_len, MPI_DOUBLE, &col_vector);
6 MPI_Type_commit (&col_vector);
7
8 // sending upwards (same for row and col)
9 if(can_transfer('u', my_rank, cluster_len, p)){
10     int target = my_rank - cluster_len;
11     MPI_Send(&data[0], 1, row_vector, target, 1, MPI_COMM_WORLD);
12 }
13
14 // recv from below (row)
15 if(can_transfer('d', my_rank, cluster_len, p)){
16     int source = my_rank + cluster_len;
17     MPI_Status st;
18     MPI_Recv(recv_data[1], 1, row_vector, source, 1, MPI_COMM_WORLD, &st);
19 }
20
21 // recv from left
22 if(can_transfer('l', my_rank, cluster_len, p)){
23     int source = my_rank - 1;
24     MPI_Status st;
25     MPI_Recv(recv_data[2], side_len, MPI_DOUBLE, source, 4, MPI_COMM_WORLD, &st);
26 }

```

Here we need to make a single MPI\_Send() MPI\_Recv() call for transfer in a single direction. Note that while receiving a row, we can safely use the row\_vector datatype, while we need to use MPI\_DOUBLE when recv for a column is made because we want to fit it in a horizontal buffer.

#### Section 3

#### Issues faced during testing/development

Following are some mentionable bugs we faced during code development/testing:

1. During MPI\_Send()/MPI\_Recv() when only one single MPI\_DOUBLE was sent/recv, we had a problem in our implementation where our send/recv buffers were overlapping.
2. During MPI\_Pack()/MPI\_Unpack(), we were not resetting the value of position variable, thus we were getting segmentation faults.
3. During MPI\_Unpack() we realized that we were losing the data when using MPI\_Isend(), this was because the call was non-blocking and therefore we were unpacking before the buffer was filled. This could be fixed in two ways, first using MPI\_Wait() for the specific call, and then unpacking. Second, to use a blocking recv so that the unpacking buffer is filled with the sender's data. We used the latter in our implementation (in all of the send/recv calls blocking send/recv are used).

4. While using `MPI_Type_vector()`, we faced a segfault because we were receiving the `col_vector` datatype in a horizontal buffer, since the layout is preserved, the `recv` call was trying to fill the `recv` buffer out of bounds thus producing the error. We fixed this by using `MPI_DOUBLE` and receiving one of it at a time.
5. While setting up the `NodeAllocator`, we faced a package not found error, so we had to manually install `psutil`. Since at that time, the CSE cluster was not connected to the internet, we had to use this script to manually connect to the internet: [https://github.com/satendrapandeymp/Auth\\_IITK/blob/master/auth-iitk.py](https://github.com/satendrapandeymp/Auth_IITK/blob/master/auth-iitk.py).

#### Section 4

##### Observations

The following 4 plots are corresponding to the number of processes involved, i.e. 16, 36, 49, 64.

- In all of the cases, multiple `send/recv` takes the highest time. This is due to the overhead for multiple `send/recv` calls.
- Using `MPI_Type_vector` takes the smallest time because of minimum overhead of type creation. In `pack/unpack`, there is an extra overhead due to packing/unpacking, while in `send/recv`, overhead is introduced due to multiple `send/recv` calls.




Figure 1: Number of Processes = 16




Figure 2: Number of Processes = 36




Figure 3: Number of Processes = 49




Figure 4: Number of Processes = 64